

APA-L3-python

September 6, 2018

1 APA Laboratori 3 - Clustering

```
In [1]: # Uncomment to upgrade packages
        # !pip install pandas --upgrade
        # !pip install numpy --upgrade
        # !pip install scipy --upgrade
        # !pip install statsmodels --upgrade
        # !pip install scikit-learn --upgrade
        %load_ext autoreload

In [2]: %%matplotlib notebook
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from collections import Counter
from IPython.core.interactiveshell import InteractiveShell
pd.set_option('precision', 3)
InteractiveShell.ast_node_interactivity = "all"

In [3]: # Extra imports
from numpy.random import uniform, normal
from sklearn.cluster import KMeans
from sklearn.metrics import calinski_harabaz_score
from numpy.random import multivariate_normal
from sklearn.datasets import make_blobs
from sklearn.mixture import GaussianMixture

In [4]: np.random.seed(7)
```

1.1 Example 1. Clustering easy artificial 2D data with k-means

First we create a simple data set:

```
In [5]: N1 = 30
        N2 = 40
        N3 = 50
```

create cluster 1

```
In [6]: x1 = normal (1,0.5,N1)
        y1 = normal (1,0.5,N1)
```

create cluster 2

```
In [7]: x2 = normal (2,0.5,N2)
        y2 = normal (6,0.7,N2)
```

create cluster 3


```
In [8]: x3 = normal (7,1,N3)
        y3 = normal (7,1,N3)
```

create the data

```
In [9]: x = np.concatenate((x1,x2,x3))
        y = np.concatenate((y1,y2,y3))
        c = np.array([0] * N1 + [1] * N2 + [2] * N3)
        D = pd.DataFrame({'X1':x, 'Y1':y, 'C':c})
```


this is your data

```
In [10]: D.plot.scatter(x='X1', y='Y1',figsize=(6,6));
```



and these are the true clusters

```
In [11]: D.plot.scatter(x='X1',  
                        y='Y1',  
                        c='C',  
                        colormap='plasma',  
                        figsize=(6,6),colorbar=False);
```



so we have 3 very clean clusters ...
Let's execute k-means

In [12]: `K = 3` # *yeah, this is tricky, why 3?*

execute k-means with a maximum of 100 iterations

In [13]: `kmeans_3 = KMeans(n_clusters=K,max_iter=100)`
`kmeans_3.fit(D.loc[:,['X1','Y1']]);`


plot and paint the clusters (according to the computed assignments) and plot the cluster centers

In [14]: `D.plot.scatter(x='X1',`
`y='Y1',`
`c=kmeans_3.labels_,`
`colormap='plasma',`
`figsize=(6,6),`

```

        colorbar=False)
plt.plot(kmeans_3.cluster_centers_[0,0],
        kmeans_3.cluster_centers_[0,1],
        'ko', markersize=12);

```



clustering quality as measured by the Calinski-Harabasz index (recommended)

This index measures the dispersion of the data points within the clusters (SSW) and between the clusters (SSB)

A good clustering has small SSW (compact clusters) and large SSB (separated cluster centers)

There is also a correction for the number of clusters

The CH index is then:

$$CH = \frac{SSB/(K-1)}{SSW/(N-K)}$$

where N is the number of data points and K is the number of clusters

```

In [15]: CH_3 = calinski_harabaz_score(D.loc[:,['X1', 'Y1']],
                                         kmeans_3.labels_ )

CH_3

```

Out[15]: 728.0901256606564

now let's not be tricky ##


In [16]: K = 5 # *guess what is going to happen?*

execute k-means with a maximum of 100 iterations

```
In [17]: kmeans_5 = KMeans(n_clusters=K,max_iter=100)
         kmeans_5.fit(D.loc[:,['X1','Y1']]);
```

plot and paint the clusters (according to the computed assignments), plot the cluster centers

```
In [18]: D.plot.scatter(x='X1',
                        y='Y1',
                        c=kmeans_5.labels_,
                        colormap='plasma',
                        figsize=(6,6),
                        colorbar=False)
plt.plot(kmeans_5.cluster_centers_[0],
         kmeans_5.cluster_centers_[1],
         'ko', markersize=12);
```



clustering quality as measured by the Calinski-Harabasz index

```
In [19]: CH_5 = calinski_harabaz_score(D.loc[:,['X1', 'Y1']],  
                                       kmeans_5.labels_)  
CH_5
```

```
Out[19]: 653.0660770798015
```

notice $CH_3 > CH_5$, so $K=3$ is better according to C-H

1.2 Example 2. Clustering not-so-easy artificial 2D data with k-means and E-M

```
In [20]: np.random.seed(6)  
K = 5  
N = 2000  
center = np.array((0,0))  
dispersion = 30  
  
mu_k = multivariate_normal(center, np.eye(2)*dispersion, K)  
pi_k = uniform(0.2, 1.5, size=K)  
data, labels = make_blobs(n_samples=N,  
                           n_features=2,  
                           centers=mu_k,  
                           cluster_std=pi_k)  
  
# Rotate and scale  
data = np.dot(data, np.array([[1, 0.2], [0.6, 1.8]]))  
# Rotate and scale  
mu_k = np.dot(mu_k, np.array([[1, 0.2], [0.6, 1.8]]))  
  
data = np.vstack((data[labels==0][:400],  
                  data[labels==1][:300],  
                  data[labels==2][:125],  
                  data[labels==3][:100],  
                  data[labels==4][:75]))  
labels = np.array([0]*400 + [1]*300 + [2]*125 + [3]*100 + [4]*75)  
d = pd.DataFrame({'X1': data[:, 0],  
                  'Y1': data[:, 1],  
                  'labels': labels})
```


generate 2D data as a mixture of 5 Gaussians, The rotation makes the two variable non independent with different variances the centers and coefficients of the mixture are chosen randomly

```
In [21]: d.plot.scatter(x='X1',  
                        y='Y1',  
                        c=labels,
```

```

        colormap='plasma',
        figsize=(6,6),
        colorbar=False)
plt.plot(mu_k[:,0], mu_k[:,1],
        'r^', markersize=12);

```



may be we want to have a look at the unconditional density $p(x)$
 compute 2D kernel density
 this is the raw data (what the clustering method sees) and a contour plot of the unconditional density

```


In [22]: d.plot.scatter(x='X1',
                        y='Y1',
                        c=labels,
                        colormap='plasma',
                        figsize=(8,8),colorbar=False)
plt.plot(mu_k[:,0], mu_k[:,1],
        'r^', markersize=12);

```




```
sns.kdeplot(d.X1,d.Y1,
            n_levels=50,cmap='afmhot');
```

```
/usr/lib64/python3.6/site-packages/scipy/stats/stats.py:1706: FutureWarning: Using a non-tuple s
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```



a simpler way of plotting the data


```
In [23]: d.plot.scatter(x='X1', y='Y1', figsize=(8,8));
```



let us try first with k-means (K=2)

In [24]: K=2

```
kmeans_2 = KMeans(n_clusters=K,max_iter=100)
kmeans_2.fit(d.loc[:,['X1','Y1']]);
d.plot.scatter(x='X1', y='Y1',
               c=kmeans_2.labels_,
               colormap='plasma',
               figsize=(6,6),colorbar=False);
```



Can we be indulgent with the result? we know the truth is there are 5 clusters,
Is this is a reasonable result if we ask for 2?
clustering quality as measured by the Calinski-Harabasz index

```
In [25]: CH_2 = calinski_harabaz_score(d.loc[:,['X1', 'Y1']],
                                         kmeans_2.labels_)
         CH_2
```


```
Out[25]: 2210.1921650543345
```

let us try now with k-means (K=5)

```
In [26]: K=5

         kmeans2_5 = KMeans(n_clusters=K,max_iter=100)
         kmeans2_5.fit(d.loc[:,['X1', 'Y1']]);
         d.plot.scatter(x='X1', y='Y1',
                        c=kmeans2_5.labels_,
```

```
colormap='plasma',
figsize=(6,6),colorbar=False);
```



This time the result has even more chances of being largely incorrect because there are more ways of getting a wrong solution
clustering quality as measured by the Calinski-Harabasz index

```
In [27]: CH2_5 = calinski_harabaz_score(d.loc[:,['X1','Y1']],
                                         kmeans2_5.labels_)

CH2_5
```

```
Out[27]: 4511.6033865718355
```

at least CH2.5 >> CH2.2 ... so C-H does a good job
In class we saw that k-means is usually re-run several times

```
In [28]: kmeans2_5 = KMeans(n_clusters=K,max_iter=100,n_init=100)
kmeans2_5.fit(d.loc[:,['X1','Y1']])
```

```
CH2_5 = calinski_harabaz_score(d.loc[:,['X1','Y1']],
                               kmeans2_5.labels_ )
CH2_5
```


```
Out[28]: KMeans(algorithm='auto', copy_x=True, init='k-means++', max_iter=100,
               n_clusters=5, n_init=100, n_jobs=1, precompute_distances='auto',
               random_state=None, tol=0.0001, verbose=0)
```

```
Out[28]: 4511.6033865718355
```

so it is not a matter of wrong initialization this is really the best k-means can do here
this may take a while

```
In [29]: r = []
         for i in range(2,11):
             km = KMeans(n_clusters=i,
                         max_iter=100,
                         n_init=100).fit(d.loc[:,['X1','Y1']])
             r.append(calinski_harabaz_score(d.loc[:,['X1','Y1']],
                                             km.labels_ ))

         fig, ax = plt.subplots(figsize=(8,6))
         plt.plot(range(2,11),r);
```



the conclusion is that k-means + C-H bet for 5 clusters ... not bad, not bad ...

but the real *shape* of the clusters cannot be captured, because k-means only "sees" spherical clusters and these are ellipsoidal

let us try now E-M

This method performs E-M for mixture densities, including mixtures of Gaussians we can specify which family of gaussians we intend to fit:

- "full" each component has full covariance matrix, "diagonal" covariances are diagonal (just variance),
- "spherical" each component has his own diagonal univariate and "tied" all components share variance

WARNING: default is "full".

suppose first that we know the truth and specify axis-aligned densities (i.e., independent variables)

```
In [30]: gm = GaussianMixture(n_components=5,
                               covariance_type='diag').fit(d.loc[:,['X1', 'Y1']])
print('BIC=', gm.bic(d.loc[:,['X1', 'Y1']]))
print('\nLOG Likelihood=', gm.lower_bound_)
print('\nWEIGHTS=')
pd.DataFrame(gm.weights_)
print('\nMEANS=')
pd.DataFrame(gm.means_)
print('\nCOV=')
pd.DataFrame(gm.covariances_)
```

BIC= 9962.697463368639

LOG Likelihood= -4.898688994806308

WEIGHTS=

```
Out [30]:      0
0  0.400
1  0.295
2  0.125
3  0.074
4  0.106
```

MEANS=

```
Out [30]:      0      1
0  0.692  6.890
1 -1.659 -8.483
```

```

2 -10.535  6.444
3  7.664 -2.046
4  0.854 -13.396

```

COV=

```

Out[30]:
      0      1
0  0.898  2.310
1  1.126  2.819
2  1.249  4.071
3  2.490  4.672
4  4.031  7.404


```

This is a graphical summary of the clustering

```

In [31]: d.plot.scatter(x='X1', y='Y1',
                        c=gm.predict(d.loc[:,['X1','Y1']]),
                        colormap='plasma',
                        figsize=(6,6),colorbar=False);

```



it was very likely that E-M performed extremely well why? because we knew the truth (cluster form and number)

suppose now we do not know the truth but we still wish to fit general gaussians

```
In [32]: gm = GaussianMixture(n_components=5,
                                covariance_type='full').fit(d.loc[:,['X1', 'Y1']])
print('BIC=', gm.bic(d.loc[:,['X1', 'Y1']]))
print('\nLOG Likelihood=', gm.lower_bound_)
print('\nWEIGHTS=')
pd.DataFrame(gm.weights_)
print('\nMEANS=')
pd.DataFrame(gm.means_)
print('\nCOV=')
gm.covariances_
```

BIC= 9609.070993641648

LOG Likelihood= -4.704401604690396

WEIGHTS=

```
Out [32]:      0
0  0.125
1  0.303
2  0.400
3  0.098
4  0.074
```

MEANS=

```
Out [32]:      0      1
0 -10.535   6.444
1  -1.692  -8.528
2   0.692   6.890
3   1.186 -13.669
4   7.667  -2.039
```

COV=

```
Out [32]: array([[1.24862465, 1.29165545],
                  [1.29165545, 4.07145078]],
```



```

[[1.13364653, 1.07469376],
 [1.07469376, 3.01171371]],

[[0.8979744 , 0.84783988],
 [0.84783988, 2.30951699]],

[[3.11324278, 2.83240408],
 [2.83240408, 6.49679973]],


[[2.49337285, 1.55130335],
 [1.55130335, 4.66268719]]])

```

```

In [33]: d.plot.scatter(x='X1', y='Y1',
                        c=gm.predict(d.loc[:,['X1','Y1']]),
                        colormap='plasma',
                        figsize=(6,6),colorbar=False);

```



the method works also very smoothly why? because the data *is* gaussian
compare the estimated centers

```
In [34]: pd.DataFrame(gm.means_)
```

```
Out[34]:
```

	0	1
0	-10.535	6.444
1	-1.692	-8.528
2	0.692	6.890
3	1.186	-13.669
4	7.667	-2.039

with the truth (note the clusters may appear in a different order)

```
In [35]: pd.DataFrame(mu_k)
```

```
Out[35]:
```

	0	1
0	0.688	6.846
1	-1.762	-8.626
2	-10.619	6.280
3	1.197	-13.693
4	7.566	-2.443

or the estimated coefficients

```
In [36]: pd.DataFrame(gm.weights_)
```

```
Out[36]:
```

	0
0	0.125
1	0.303
2	0.400
3	0.098
4	0.074

with the truth

```
In [37]: c= Counter(labels)
          pd.DataFrame(np.array([c[v] for v in c])/1000.0)
```

```
Out[37]:
```

	0
0	0.400
1	0.300
2	0.125
3	0.100
4	0.075