

# APA-L7-python

September 6, 2018

## 1 APA Laboratori 7 - SVMs

```
In [1]: # Uncomment to upgrade packages
        # !pip install pandas --upgrade
        # !pip install numpy --upgrade
        # !pip install scipy --upgrade
        # !pip install statsmodels --upgrade
        # !pip install scikit-learn --upgrade
        %load_ext autoreload

In [2]: %%matplotlib notebook
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import matplotlib.cm as cm
from IPython.core.interactiveshell import InteractiveShell
pd.set_option('precision', 3)
InteractiveShell.ast_node_interactivity = "all"

In [3]: # Extra imports
from numpy.random import uniform, normal
from sklearn.svm import SVC
from sklearn.model_selection import KFold
from sklearn.metrics import confusion_matrix, accuracy_score
from sklearn.svm import SVR

In [4]: np.random.seed(7)
```

### 1.1 Modelling artificial 2D sinusoidal data for two-class problems

First we create a simple two-class data set:

```
In [5]: N = 200

def make_sinusoidals(m, noise=0.2):
    x1 = np.ones(2*m)
    x2 = np.ones(2*m)
```

```

for i in range(m):
    x1[i] = (i/m) * np.pi
    x2[i] = np.sin(x1[i]) + normal(0,noise,1)

for j in range(m):
    x1[m+j] = (j/m + 1/2) * np.pi
    x2[m+j] = np.cos(x1[m+j]) + normal(0,noise,1)

target = [1]*m+[-1]*m

return pd.DataFrame({'x1':x1,'x2':x2,'t':target})

```

let's generate the data

```
In [6]: dataset = make_sinusoidals(N)
```


and have a look at it

```
In [7]: dataset.describe()
```

```
Out[7]:
```

	x1	x2	t
count	400.000	400.000	400.000
mean	2.348	-0.010	0.000
std	1.201	0.739	1.001
min	0.000	-1.402	-1.000
25%	1.567	-0.675	-1.000
50%	2.348	-0.040	0.000
75%	3.130	0.677	1.000
max	4.697	1.450	1.000

```
In [8]: dataset.plot.scatter(x='x1', y='x2', c='t',
                             colormap='plasma',
                             figsize=(8,8),
                             colorbar=False);
```



Now we wish to fit and visualize different SVM models

### 1.1.1 model 1: LINEAR kernel, C=1 (cost parameter)

```
In [9]: svc = SVC(C=1, kernel='linear')
        svc.fit(dataset.loc[:, 'x1': 'x2'], dataset.t)
```

```
print('Suports=', svc.n_support_)
```

```
Out[9]: SVC(C=1, cache_size=200, class_weight=None, coef0=0.0,
            decision_function_shape='ovr', degree=3, gamma='auto', kernel='linear',
            max_iter=-1, probability=False, random_state=None, shrinking=True,
            tol=0.001, verbose=False)
```

```
Suports= [26 25]
```

Now we are going to visualize what we have done; since we have artificial data, instead of creating a random test set, we can create a grid of points as test


```
In [10]: def plot_prediction(model, model_name, resol=200, ax=None):
    x_min, x_max = dataset.x1.min(), dataset.x1.max()
    y_min, y_max = dataset.x2.min(), dataset.x2.max()
    xx, yy = np.meshgrid(np.linspace(x_min, x_max, resol),
                          np.linspace(y_min, y_max, resol))

    Z = model.decision_function(np.c_[xx.ravel(), yy.ravel()])
    Z = Z.reshape(xx.shape)
    if ax is None:
        fig, ax = plt.subplots(figsize=(8,8))
    ax.imshow(Z, interpolation='bilinear',
              cmap=cm.RdYlGn,
              extent=[x_min, x_max, y_min, y_max])
    dataset.plot.scatter(x='x1', y='x2', c='t',
                        colormap='plasma',
                        colorbar=False,
                        ax=ax,
                        title=model_name+', supp=%d'%np.sum(model.n_support_))
    dataset.iloc[model.support_].plot.scatter(x='x1', y='x2', c='t',
                                              colormap='bwr',
                                              colorbar=False,
                                              ax=ax, marker='x',
                                              s=40)

    CS=ax.contour(xx, yy, Z, levels=[-1,0,1], colors=['grey', 'black', 'grey'])
    plt.clabel(CS, inline=1, fontsize=10)
```

make sure you understand the following results (one by one and their differences)  
plot the data, the OSH with margins, the support vectors, ...

```
In [11]: plot_prediction(svc, 'linear, C=1')
```



### model 2: linear kernel, C=0.1 (cost parameter)


```
In [12]: svc = SVC(C=0.1, kernel='linear')
         svc.fit(dataset.loc[:, 'x1': 'x2'], dataset.t)
```

```
print('Suports=', svc.n_support_)
```

```
Out[12]: SVC(C=0.1, cache_size=200, class_weight=None, coef0=0.0,
             decision_function_shape='ovr', degree=3, gamma='auto', kernel='linear',
             max_iter=-1, probability=False, random_state=None, shrinking=True,
             tol=0.001, verbose=False)
```

```
Suports= [49 48]
```

```
In [13]: plot_prediction(svc, 'linear, C=0.1')
```



the margin is wider (lower VC dimension), number of support vectors is larger (more violations of the margin)

### 1.1.2 model 3: linear kernel, C=25 (cost parameter)

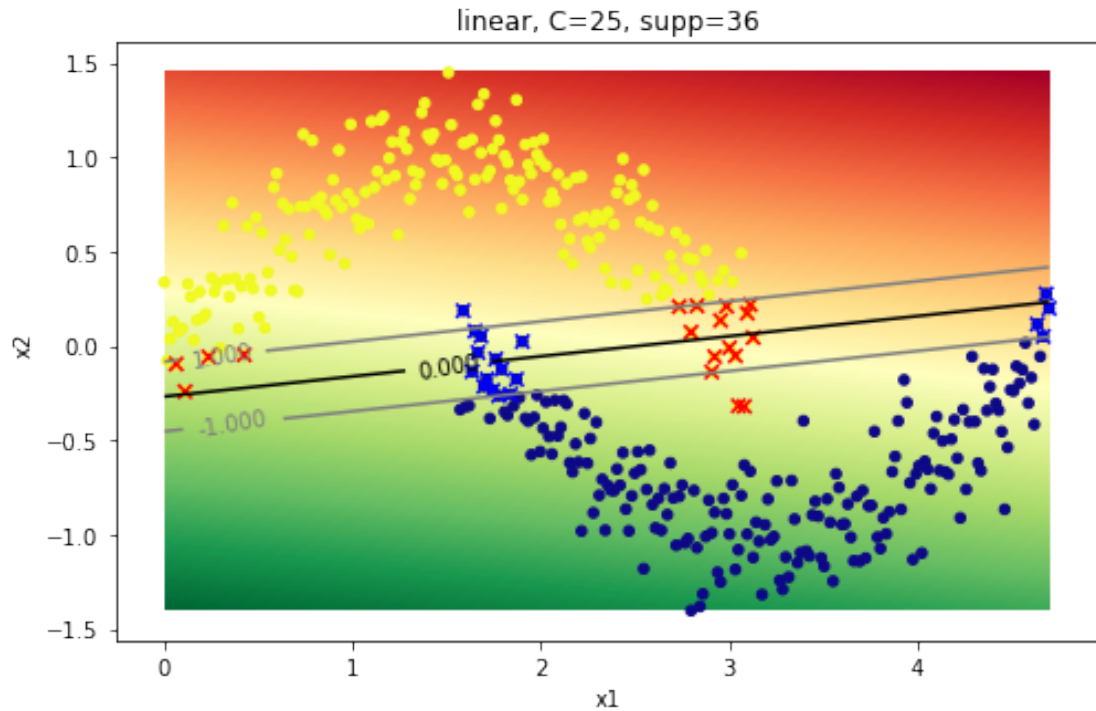
```
In [14]: svc = SVC(C=25, kernel='linear')
         svc.fit(dataset.loc[:, 'x1': 'x2'], dataset.t)
```

```
print('Suports=', svc.n_support_)
```

```
Out[14]: SVC(C=25, cache_size=200, class_weight=None, coef0=0.0,
             decision_function_shape='ovr', degree=3, gamma='auto', kernel='linear',
             max_iter=-1, probability=False, random_state=None, shrinking=True,
             tol=0.001, verbose=False)
```

```
Suports= [18 18]
```

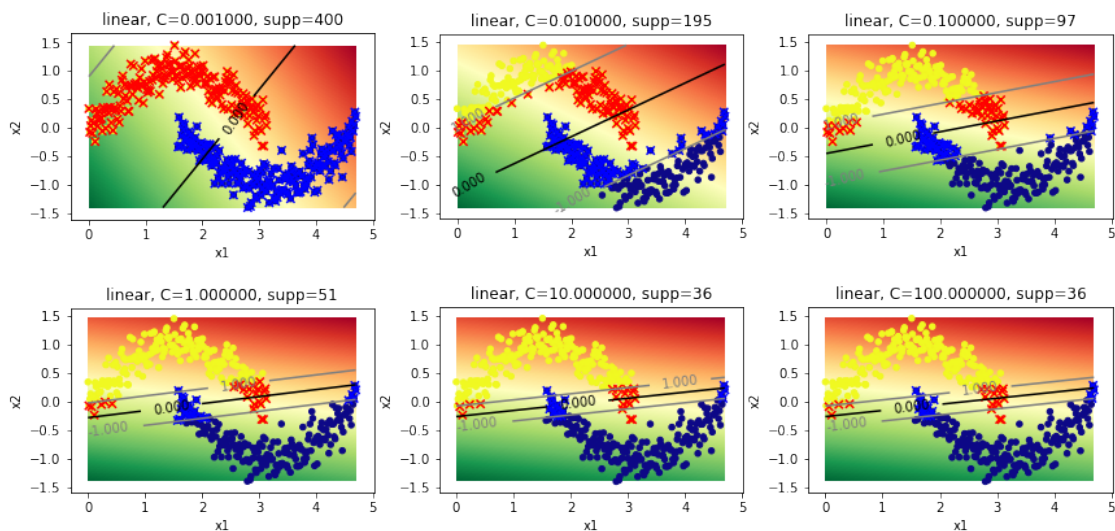
```
In [15]: plot_prediction(svc, 'linear, C=25')
```



the margin is narrower (higher VC dimension), number of support vectors is smaller (less violations of the margin)

Let's put it together, for 6 values of C:

```
In [16]: fig, ax = plt.subplots(2,3,figsize=(12,6))
         for C, ax in zip(10**np.linspace(-3,2,6),ax.ravel()):
             svc = SVC(C=C, kernel='linear').fit(dataset.loc[:, 'x1': 'x2'], dataset.t)
             plot_prediction(svc, 'linear, C=%f'%C, ax=ax)
         plt.tight_layout()
```



Now we move to a QUADRATIC kernel (polynomial of degree 2); the kernel has the form:

$$k(x, y) = (\langle x, y \rangle + \text{coef0})^{\text{degree}}$$

quadratic kernel, C=1 (cost parameter)


```
In [17]: svc = SVC(C=1, kernel='poly', degree=2, coef0=1)
        svc.fit(dataset.loc[:, 'x1': 'x2'], dataset.t)
```

```
print('Suports=', svc.n_support_)
```

```
Out[17]: SVC(C=1, cache_size=200, class_weight=None, coef0=1,
            decision_function_shape='ovr', degree=2, gamma='auto', kernel='poly',
            max_iter=-1, probability=False, random_state=None, shrinking=True,
            tol=0.001, verbose=False)
```

Suports= [21 22]

```
In [18]: plot_prediction(svc, 'quadratic, C=1')
```




notice that neither the OSH or the margins are linear (they are quadratic); they are linear in the feature space in the previous linear kernel, both spaces coincide

Let's put it together directly, for 6 values of C:



```
In [19]: fig, ax = plt.subplots(2,3,figsize=(12,6))
         for C, ax in zip(10**np.linspace(-3,2,6),ax.ravel()):
             svc = SVC(C=C,
                        kernel='poly',
                        degree=2,
                        coef0=1).fit(dataset.loc[:, 'x1': 'x2'],
                                    dataset.t)
             plot_prediction(svc, 'quadratic, C=%f'%C, ax=ax)
         plt.tight_layout()
```



Now we move to a CUBIC kernel (polynomial of degree 3); the kernel has the form:

$$k(x, y) = (\langle x, y \rangle + \text{coef0})^{\text{degree}}$$

cubic kernel, C=1 (cost parameter)

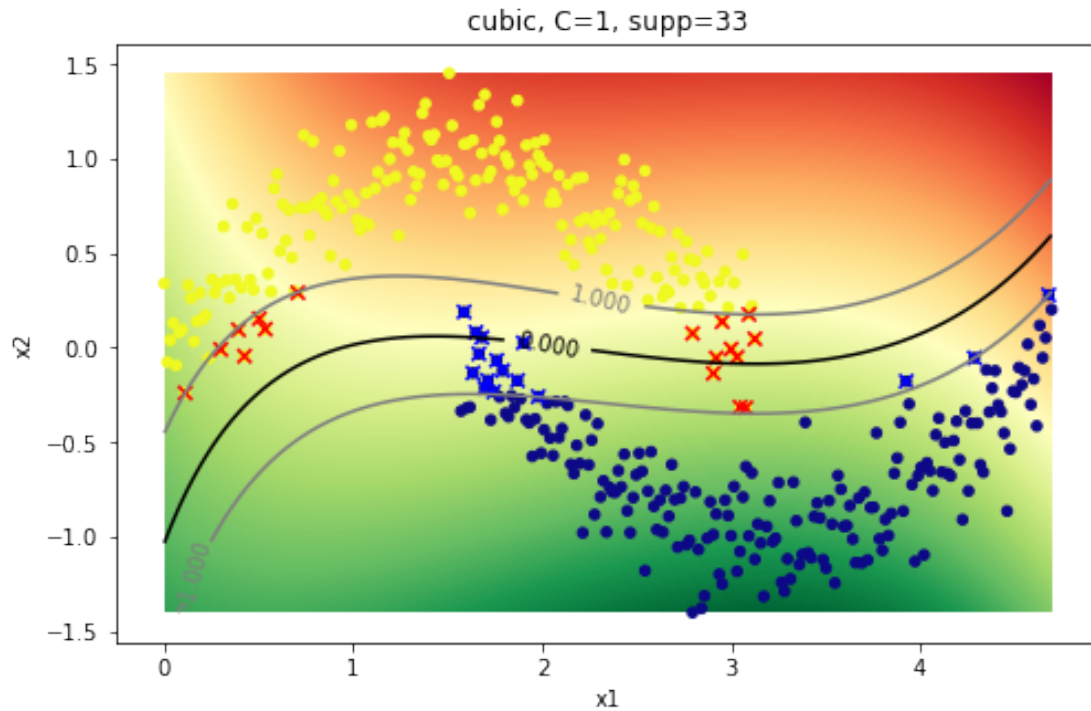
```
In [20]: svc = SVC(C=1, kernel='poly', degree=3, coef0=1)
         svc.fit(dataset.loc[:, 'x1': 'x2'], dataset.t)

         print('Suports=', svc.n_support_)

Out[20]: SVC(C=1, cache_size=200, class_weight=None, coef0=1,
             decision_function_shape='ovr', degree=3, gamma='auto', kernel='poly',
             max_iter=-1, probability=False, random_state=None, shrinking=True,
             tol=0.001, verbose=False)
```

Suports= [16 17]

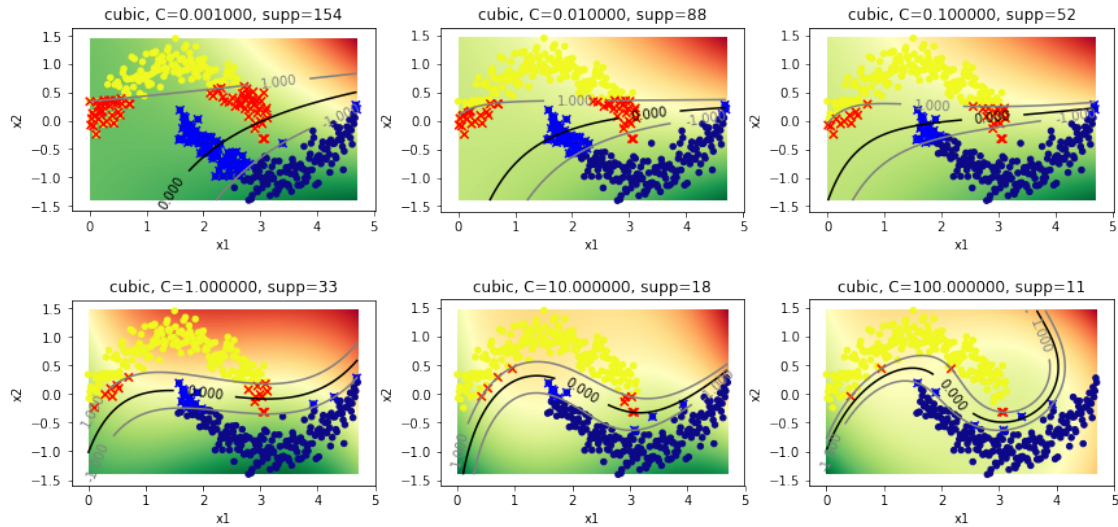
```
In [21]: plot_prediction(svc, 'cubic, C=1')
```



notice that neither the OSH or the margins are linear (they are now cubic); they are linear in the feature space this choice seems much better, given the structure of the classes

Let's put it together directly, for 6 values of C:

```
In [22]: fig, ax = plt.subplots(2,3,figsize=(12,6))
         for C, ax in zip(10*np.linspace(-3,2,6),ax.ravel()):
             svc = SVC(C=C,
                        kernel='poly',
                        degree=3,
                        coef0=1).fit(dataset.loc[:, 'x1': 'x2'],
                                    dataset.t)
             plot_prediction(svc, 'cubic, C=%f'%C,ax=ax)
         plt.tight_layout()
```



Finally we use the Gaussian RBF kernel (polynomial of infinite degree; the kernel has the form:

$$k(x, y) = \exp(-\gamma \|x - y\|^2)$$

RBF kernel, C=1 (cost parameter)

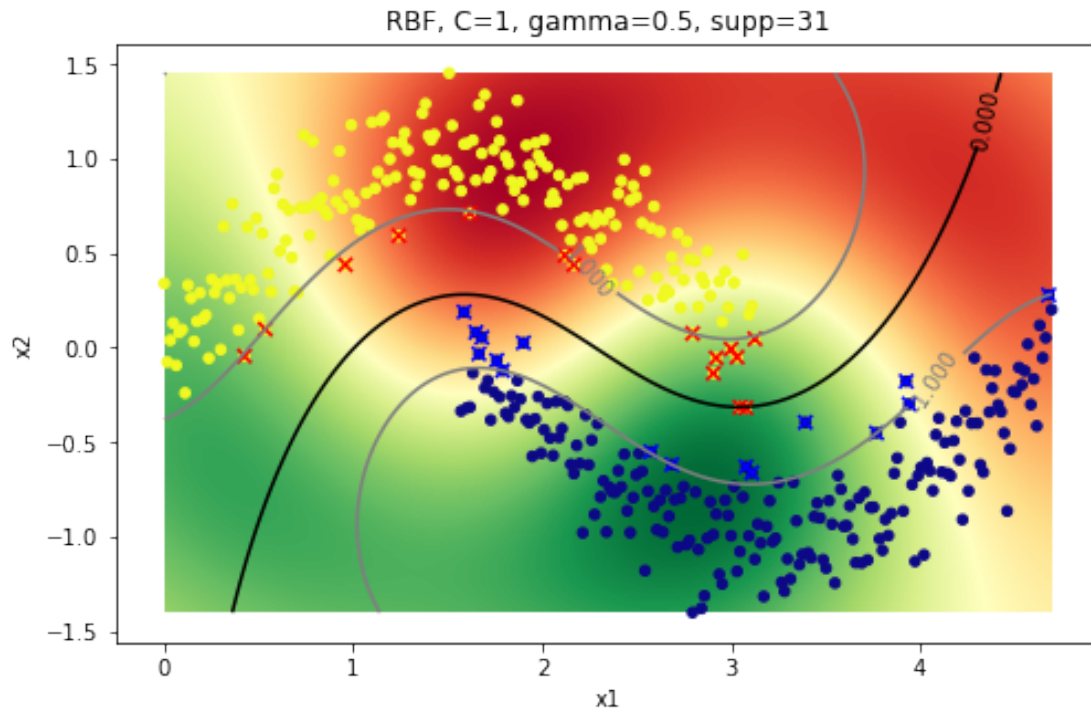
```
In [23]: svc = SVC(C=1, kernel='rbf', gamma=0.5)
         svc.fit(dataset.loc[:, 'x1': 'x2'], dataset.t)

         print('Suports=', svc.n_support_)
```

```
Out[23]: SVC(C=1, cache_size=200, class_weight=None, coef0=0.0,
            decision_function_shape='ovr', degree=3, gamma=0.5, kernel='rbf',
            max_iter=-1, probability=False, random_state=None, shrinking=True,
            tol=0.001, verbose=False)
```

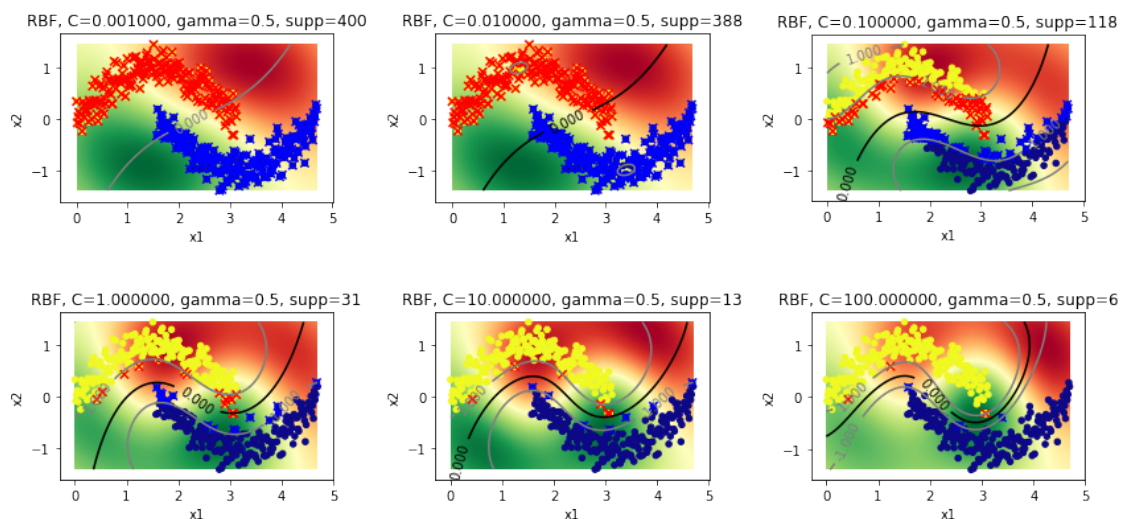
```
Suports= [16 15]
```

```
In [24]: plot_prediction(svc, 'RBF, C=1, gamma=0.5')
```



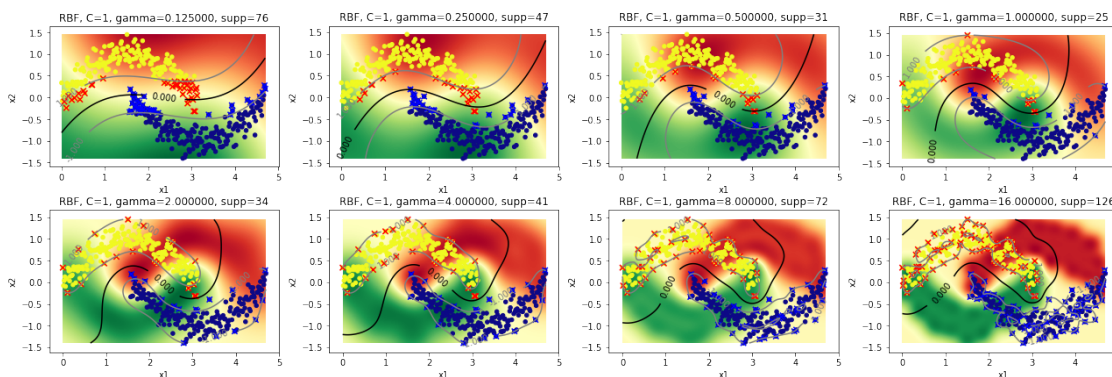
Let's put it together directly, for 6 values of C, holding gamma constant = 0.5:

```
In [25]: fig, ax = plt.subplots(2,3,figsize=(12,6))
        for C, ax in zip(10**np.linspace(-3,2,6),ax.ravel()):
            svc = SVC(C=C,
                      kernel='rbf',
                      gamma=0.5).fit(dataset.loc[:, 'x1': 'x2'],
                                      dataset.t)
            plot_prediction(svc, 'RBF, C=%f, gamma=0.5'%C, ax=ax)
        plt.tight_layout()
```



Now for 8 values of gamma, holding C constant = 1:

```
In [26]: fig, ax = plt.subplots(2,4,figsize=(18,6))
         for gamma, ax in zip(2**np.linspace(-3,4,8),ax.ravel()):
             svc = SVC(C=1,
                        kernel='rbf',
                        gamma=gamma).fit(dataset.loc[:, 'x1':'x2'],
                                       dataset.t)
             plot_prediction(svc, 'RBF, C=1, gamma=%f'%gamma,ax=ax)
         plt.tight_layout()
```



In practice we should optimize both (C,gamma) at the same time

How? Using cross-validation or trying to get "good" estimates analyzing the data

Now we define a utility function for performing k-fold CV: a typical choice is k=10

```
In [27]: k = 10
```

```
In [28]: def train_svm_kCV(C, k, kCV=10, params={'kernel':'linear'}):
         merr = 0.0
         for i in range(kCV):
             cv = KFold(n_splits=k)
             for tr_ind,ts_ind in cv.split(dataset):
                 svc = SVC(C=C,
                           **params).fit(dataset.iloc[tr_ind,[0,1]],
                                         dataset.t.iloc[tr_ind])
                 pred = svc.predict(dataset.iloc[ts_ind,[0,1]])
                 merr += (1-accuracy_score(dataset.t.iloc[ts_ind],pred))
         return merr/(kCV*k)*100
```

```
In [29]: C = 10
```

```
VA_error_linear = train_svm_kCV (C=C, k=10, params={'kernel':'linear'})
VA_error_linear
```

```
Out [29]: 17.749999999999996
```

procedure is to choose the model with the lowest CV error and then refit it with the whole learning data, then use it to predict the test set; we will do this at the end

Fit an SVM with quadratic kernel

```
In [30]: VA_error_poly_2 = train_svm_kCV (C=C, k=10,  
                                         params={'kernel':'poly', 'degree':2,'coef0':1})  
VA_error_poly_2
```

```
Out [30]: 25.250000000000007
```

Fit an SVM with cubic kernel

```
In [31]: VA_error_poly_3 = train_svm_kCV (C=C, k=10,  
                                         params={'kernel':'poly', 'degree':3,'coef0':1})  
VA_error_poly_3
```

```
Out [31]: 6.7500000000000036
```

```
In [32]: VA_error_rbf = train_svm_kCV (C=C, k=10, params={'kernel':'rbf'})  
VA_error_rbf
```

```
Out [32]: 1.25
```

Now in a real scenario we should choose the model with the lowest CV error which in this case is the RBF (we get a very low CV error because this problem is easy for a SVM)  
so we choose RBF and C=1 and refit the model in the whole training set (no CV)

```
In [33]: svc = SVC(C=1, kernel='rbf').fit(dataset.loc[:, 'x1': 'x2'], dataset.t)
```

and make it predict a test set:  
let's generate the test data

```
In [34]: dataset_test = make_sinusoidals (1000)
```


and have a look at it

```
In [35]: dataset_test.describe()
```

```
Out [35]:
```

	x1	x2	t
count	2000.000	2000.000	2000.0
mean	2.355	0.005	0.0
std	1.200	0.730	1.0
min	0.000	-1.570	-1.0
25%	1.570	-0.653	-1.0
50%	2.355	0.003	0.0
75%	3.139	0.688	1.0
max	4.709	1.523	1.0

```
In [36]: dataset_test.plot.scatter(x='x1', y='x2', c='t',
                                   colormap='plasma',
                                   figsize=(8,8),
                                   colorbar=False);
```



```
In [37]: pred = svc.predict(dataset_test.loc[:, 'x1': 'x2'])
df_tr=pd.DataFrame(confusion_matrix(dataset_test.t, pred),
                    index=['-1', '1'], columns=['-1', '1'])

df_tr
(1-accuracy_score(dataset_test.t, pred))*100
```

```
Out [37]:
```

	-1	1
-1	996	4
1	5	995

```
Out [37]: 0.4499999999999999485
```



In a real setting we should also optimize the value of  $C$ , again with CV; all this can be done very conveniently using GridSearchCV to do automatic grid-search (very much as we did in the last laboratory for nnet)

other packages provide with heuristic methods to estimate the gamma in the RBF kernel (see below)

[Here](#) you can find an interactive tool that allows to play with SVM

---


## 1.2 Playing with the SVM for regression and 1D data

Now we do regression; we have an extra parameter: the 'epsilon', which controls the width of the epsilon-insensitive tube (in feature space)

In [38]:  $A=20$

a really nice-looking function


```
In [39]: x = np.arange(-A,A,0.11)
y = (np.sin(x)/x) + normal(scale=0.03,size=len(x))
x = np.array(x).reshape(-1, 1)
fig, ax = plt.subplots(figsize=(10,6))
plt.plot(x,y, 'k-');
```



With this choice of the 'epsilon', 'gamma' and  $C$  parameters, the SVM underfits the data (blue line)

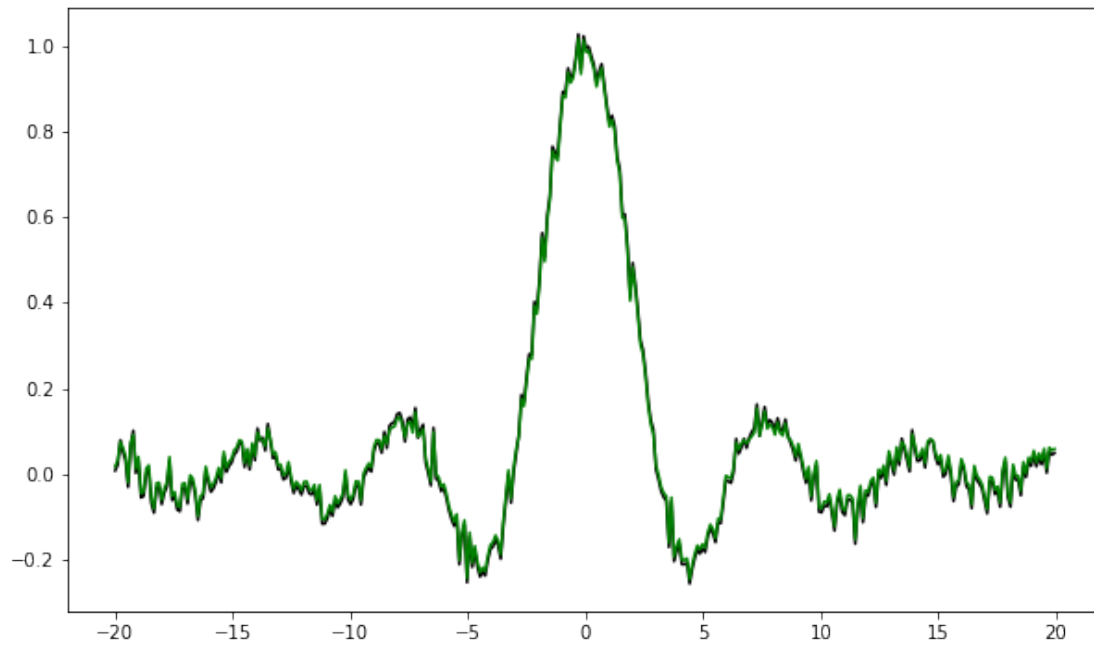


```
In [40]: svr = SVR(epsilon=0.1,gamma=0.1).fit(x,y)
fig, ax = plt.subplots(figsize=(10,6))
plt.plot(x,y, 'k-')
plt.plot(x,svr.predict(x), 'b-');
```



With this choice of the 'epsilon', 'gamma' and C parameters, the SVM overfits the data (green line)

```
In [41]: svr = SVR(epsilon=0.01,gamma=200, C=100).fit(x,y)
fig, ax = plt.subplots(figsize=(10,6))
plt.plot(x,y, 'k-')
plt.plot(x,svr.predict(x), 'g-');
```



With this choice of the 'epsilon', 'gamma' and C parameters, the SVM has a very decent fit (red line)

```
In [42]: svr = SVR(epsilon=0.01,gamma=1).fit(x,y)
fig, ax = plt.subplots(figsize=(10,6))
plt.plot(x,y,'k-')
plt.plot(x,svr.predict(x),'r-');
```

