

APA-L7

September 6, 2018

1 APA Laboratori 7 - SVMs

```
In [1]: # uncomment to install missing libraries
        # install.packages('e1071')
```

```
In [2]: options(repr.plot.width=6, repr.plot.height=6)
```

1.1 Modelling artificial 2D sinusoidal data for two-class problems

The SVM is located in two different packages: one of them is 'e1071'

```
In [3]: library(e1071)
```

First we create a simple two-class data set:

```
In [4]: N <- 200
```

```
make.sinusoidals <- function(m,noise=0.2)
{
  x1 <- c(1:2*m)
  x2 <- c(1:2*m)

  for (i in 1:m) {
    x1[i] <- (i/m) * pi
    x2[i] <- sin(x1[i]) + rnorm(1,0,noise)
  }

  for (j in 1:m) {
    x1[m+j] <- (j/m + 1/2) * pi
    x2[m+j] <- cos(x1[m+j]) + rnorm(1,0,noise)
  }

  target <- as.factor(c(rep(+1,m),rep(-1,m)))

  return(data.frame(x1,x2,target))
}
```

let's generate the data


```
In [5]: dataset <- make.sinusoidals (N)
```

and have a look at it

```
In [6]: summary(dataset)
```

	x1	x2	target
Min.	:0.01571	Min. :-1.387200	-1:200
1st Qu.:	1.58258	1st Qu.: -0.668201	1 :200
Median :	2.36405	Median : 0.028719	
Mean :	2.36405	Mean :-0.006386	
3rd Qu.:	3.14552	3rd Qu.: 0.645680	
Max.	:4.71239	Max. : 1.386127	

```
In [7]: options(repr.plot.width=4, repr.plot.height=4)
plot(dataset$x1,dataset$x2,col=dataset$target)
```



Now we wish to fit and visualize different SVM models

1.1.1 model 1: LINEAR kernel, C=1 (cost parameter)

```
In [8]: (model <- svm(dataset[,1:2],
                      dataset[,3],
```

```

type="C-classification",
cost=1, kernel="linear",
scale = FALSE))

```

Call:

```

svm.default(x = dataset[, 1:2], y = dataset[, 3], scale = FALSE,
  type = "C-classification", kernel = "linear", cost = 1)

```

Parameters:

```

SVM-Type: C-classification
SVM-Kernel: linear
cost: 1
gamma: 0.5

```

Number of Support Vectors: 61

Now we are going to visualize what we have done; since we have artificial data, instead of creating a random test set, we can create a grid of points as test

```

In [9]: plot.prediction <- function (model, model.name, resol=200)
  # the grid has a (resol x resol) resolution
  {
    x <- cbind(dataset$x1,dataset$x2)
    rng <- apply(x,2,range);
    tx <- seq(rng[1,1],rng[2,1],length=resol);
    ty <- seq(rng[1,2],rng[2,2],length=resol);
    pnts <- matrix(nrow=length(tx)*length(ty),ncol=2);
    k <- 1
    for(j in 1:length(ty))
    {
      for(i in 1:length(tx))
      {
        pnts[k,] <- c(tx[i],ty[j])
        k <- k+1
      }
    }

    # we calculate the predictions on the grid

    pred <- predict(model, pnts, decision.values = TRUE)

    z <- matrix(attr(pred,"decision.values"),
      nrow=length(tx),ncol=length(ty))
  }

```

```

# and plot them

image(tx,ty,z,xlab=model.name,ylab="",axes=FALSE,
      xlim=c(rng[1,1],rng[2,1]),ylim=c(rng[1,2],rng[2,2]),
      col = rainbow(200, start=0, end=.25))

# then we draw the optimal separation and its margins

contour(tx,ty,z,add=TRUE,
        drawlabels=TRUE,
        level=0, lwd=3)
contour(tx,ty,z,add=TRUE,
        drawlabels=TRUE, level=1,
        lty=1, lwd=1, col="grey")
contour(tx,ty,z,add=TRUE,
        drawlabels=TRUE, level=-1,
        lty=1, lwd=1, col="grey")

# then we plot the input data from the two classes

points(dataset[dataset$target==1,1:2],pch=21,col=1,cex=1)
points(dataset[dataset$target==-1,1:2],pch=19,col=4,cex=1)


# finally we add the SVs

sv <- dataset[c(model$index),];
sv1 <- sv[sv$target==1,];
sv2 <- sv[sv$target==-1,];
points(sv1[,1:2],pch=13,col=1,cex=2)
points(sv2[,1:2],pch=13,col=4,cex=2)
}

```

make sure you understand the following results (one by one and their differences)
 plot the data, the OSH with margins, the support vectors, ...

In [10]: `plot.prediction (model, "linear, C=1")`



linear, C=1

model 2: linear kernel, C=0.1 (cost parameter)

```
In [11]: (model <- svm(dataset[,1:2],dataset[,3],
                      type="C-classification",
                      cost=0.1, kernel="linear",
                      scale = FALSE))
```

Call:


```
svm.default(x = dataset[, 1:2], y = dataset[, 3], scale = FALSE,
            type = "C-classification", kernel = "linear", cost = 0.1)
```

Parameters:

```
SVM-Type: C-classification
SVM-Kernel: linear
cost: 0.1
gamma: 0.5
```

Number of Support Vectors: 99

```
In [12]: plot.prediction (model, "linear, C=0.1")
```



linear, $C=0.1$

the margin is wider (lower VC dimension), number of support vectors is larger (more violations of the margin)

1.1.2 model 3: linear kernel, $C=25$ (cost parameter)

```
In [13]: (model <- svm(dataset[,1:2],dataset[,3],
                      type="C-classification",
                      cost=25, kernel="linear",
                      scale = FALSE))
```

Call:


```
svm.default(x = dataset[, 1:2], y = dataset[, 3], scale = FALSE,
           type = "C-classification", kernel = "linear", cost = 25)
```

Parameters:

```
SVM-Type: C-classification
SVM-Kernel: linear
cost: 25
gamma: 0.5
```

Number of Support Vectors: 45

```
In [14]: plot.prediction (model, "linear, C=25")
```




linear, C=25

the margin is narrower (higher VC dimension), number of support vectors is smaller (less violations of the margin)


Let's put it together, for 6 values of C:

```
In [15]: options(repr.plot.width=6, repr.plot.height=5)
par(mfrow=c(2,3))


for (C in 10^seq(-2,3))
{
  model <- svm(dataset[,1:2],dataset[,3],
               type="C-classification",
               cost=C, kernel="linear",
               scale = FALSE)
  plot.prediction (model, paste ("linear (C=", C, ") ",
                                model$tot.nSV, " Support Vectors",
                                sep=""))
}
```




linear (C=0.01) 196 Support Vectors




linear (C=0.1) 99 Support Vectors




linear (C=1) 61 Support Vectors



linear (C=10) 46 Support Vectors



linear (C=100) 44 Support Vectors



linear (C=1000) 43 Support Vectors

Now we move to a QUADRATIC kernel (polynomial of degree 2); the kernel has the form:
 $k(x, y) = (\langle x, y \rangle + \text{coef0})^{\text{degree}}$
 quadratic kernel, C=1 (cost parameter)

```
In [16]: (model <- svm(dataset[,1:2],dataset[,3],
                        type="C-classification",
                        cost=1, kernel="polynomial",
                        degree=2, coef0=1,
                        scale = FALSE))
```

Call:

```
svm.default(x = dataset[, 1:2], y = dataset[, 3], scale = FALSE,
            type = "C-classification", kernel = "polynomial", degree = 2,
            coef0 = 1, cost = 1)
```

Parameters:

SVM-Type: C-classification


```

SVM-Kernel:  polynomial
             cost:  1
             degree: 2
             gamma:  0.5
             coef.0: 1

```


Number of Support Vectors: 48

```

In [17]: options(repr.plot.width=4, repr.plot.height=4)
         par(mfrow=c(1,1))

         plot.prediction (model, "quadratic, C=1")

```



quadratic, C=1

notice that neither the OSH or the margins are linear (they are quadratic); they are linear in the feature space in the previous linear kernel, both spaces coincide

Let's put it together directly, for 6 values of C:

```

In [18]: options(repr.plot.width=6, repr.plot.height=5)
         par(mfrow=c(2,3))


         for (C in 10^seq(-2,3))
         {

```


```

model <- svm(dataset[,1:2],dataset[,3],
             type="C-classification",
             cost=C, kernel="polynomial",
             degree=2, coef0=1,
             scale = FALSE)
plot.prediction (model,
                 paste ("quadratic (C=", C, ") ",
                       model$tot.nSV, " Support Vectors",
                       sep=""))
}


```




quadratic (C=0.01) 131 Support Vectors




quadratic (C=0.1) 72 Support Vectors




quadratic (C=1) 48 Support Vectors



quadratic (C=10) 39 Support Vectors



quadratic (C=100) 38 Support Vectors



quadratic (C=1000) 38 Support Vectors

Now we move to a CUBIC kernel (polynomial of degree 3); the kernel has the form:

$$k(x,y) = (\langle x,y \rangle + \text{coef0})^{\text{degree}}$$

cubic kernel, C=1 (cost parameter)

```

In [19]: (model <- svm(dataset[,1:2],dataset[,3],
                      type="C-classification",
                      cost=1, kernel="polynomial",
                      degree=3, coef0=1,
                      scale = FALSE))

```

Call:


```
svm.default(x = dataset[, 1:2], y = dataset[, 3], scale = FALSE,  
            type = "C-classification", kernel = "polynomial", degree = 3,  
            coef0 = 1, cost = 1)
```

Parameters:

```
SVM-Type: C-classification  
SVM-Kernel: polynomial  
cost: 1  
degree: 3  
gamma: 0.5  
coef.0: 1
```

Number of Support Vectors: 37

```
In [20]: options(repr.plot.width=4, repr.plot.height=4)  
         par(mfrow=c(1,1))  
  
         plot.prediction (model, "cubic, C=1")
```




cubic, C=1

notice that neither the OSH or the margins are linear (they are now cubic); they are linear in the feature space this choice seems much better, given the structure of the classes


Let's put it together directly, for 6 values of C:

```
In [21]: options(repr.plot.width=6, repr.plot.height=5)
         par(mfrow=c(2,3))


         for (C in 10^seq(-2,3))
         {
             model <- svm(dataset[,1:2],dataset[,3],
                           type="C-classification",
                           cost=C, kernel="polynomial",
                           degree=3, coef0=1,
                           scale = FALSE)
             plot.prediction (model, paste ("cubic (C=", C, ") ",
                                             model$tot.nSV,
                                             " Support Vectors",
                                             sep=""))
         }
```




cubic (C=0.01) 88 Support Vectors




cubic (C=0.1) 54 Support Vectors




cubic (C=1) 37 Support Vectors



cubic (C=10) 16 Support Vectors



cubic (C=100) 8 Support Vectors



cubic (C=1000) 7 Support Vectors

Finally we use the Gaussian RBF kernel (polynomial of infinite degree; the kernel has the form:
 $k(x,y) = \exp(-\gamma||x - y||^2)$
RBF kernel, C=1 (cost parameter)

```
In [22]: (model <- svm(dataset[,1:2],dataset[,3],
                      type="C-classification",
                      cost=1, kernel="radial",
                      scale = FALSE))
```

Call:

```
svm.default(x = dataset[, 1:2], y = dataset[, 3], scale = FALSE,
            type = "C-classification", kernel = "radial", cost = 1)
```

Parameters:


```
SVM-Type:  C-classification
SVM-Kernel: radial
cost:      1
gamma:     0.5
```

Number of Support Vectors: 34

the default value for gamma is 0.5

```
In [23]: options(repr.plot.width=4, repr.plot.height=4)
         par(mfrow=c(1,1))

         plot.prediction (model, "radial, C=1, gamma=0.5")
```




radial, C=1, gamma=0.5


Let's put it together directly, for 6 values of C, holding gamma constant = 0.5:

```
In [24]: options(repr.plot.width=6, repr.plot.height=5)
         par(mfrow=c(2,3))


         for (C in 10^seq(-2,3))
         {
             model <- svm(dataset[,1:2],dataset[,3],
                           type="C-classification",
                           cost=C, kernel="radial",
                           scale = FALSE)
             plot.prediction (model, paste ("RBF (C=", C, ") ",
                                             model$tot.nSV,
                                             " Support Vectors",
                                             sep=""))
         }
```




RBf (C=0.01) 390 Support Vectors




RBf (C=0.1) 123 Support Vectors




RBf (C=1) 34 Support Vectors



RBf (C=10) 12 Support Vectors



RBf (C=100) 6 Support Vectors




RBf (C=1000) 5 Support Vectors


Now for 8 values of gamma, holding C constant = 1:

```
In [25]: options(repr.plot.width=8, repr.plot.height=5)
         par(mfrow=c(2,4))

         for (g in 2^seq(-3,4))
         {
             model <- svm(dataset[,1:2],dataset[,3],
                           type="C-classification",
                           cost=1, kernel="radial",
                           gamma=g, scale = FALSE)
             plot.prediction (model, paste ("RBf (gamma=", g, ") ",
                                             model$tot.nSV,
                                             " Support Vectors",
                                             sep=""))
         }
```



RBF (gamma=0.125) 80 Support Vectors RBF (gamma=0.25) 54 Support Vectors RBF (gamma=0.5) 34 Support Vectors RBF (gamma=1) 30 Support Vectors



RBF (gamma=2) 34 Support Vectors RBF (gamma=4) 37 Support Vectors RBF (gamma=8) 73 Support Vectors RBF (gamma=16) 130 Support Vectors

In practice we should optimize both (C, γ) at the same time

How? Using cross-validation or trying to get "good" estimates analyzing the data

Now we define a utility function for performing k-fold CV: a typical choice is $k=10$

```
In [26]: k <- 10
         folds <- sample(rep(1:k, length=N), N, replace=FALSE)

         valid.error <- rep(0,k)
```

this function is not intended to be useful for general training purposes but it is useful for illustration in particular, it does not optimize the value of C (it requires it as parameter)

```
In [27]: train.svm.kCV <- function (which.kernel, myC, kCV=10)
{
  for (i in 1:kCV)
  {
    train <- dataset[folds!=i,] # for building the model (training)
    valid <- dataset[folds==i,] # for prediction (validation)

    x_train <- train[,1:2]
    t_train <- train[,3]

    switch(which.kernel,
           linear={model <- svm(x_train, t_train,
                                type="C-classification",
                                cost=myC, kernel="linear",
```



```

                                scale = FALSE)),
poly.2={model <- svm(x_train, t_train,
                    type="C-classification",
                    cost=myC, kernel="polynomial",
                    degree=2, coef0=1,
                    scale = FALSE)},
poly.3={model <- svm(x_train, t_train,
                    type="C-classification",
                    cost=myC,
                    kernel="polynomial",
                    degree=3, coef0=1,
                    scale = FALSE)},
RBF={model <- svm(x_train, t_train,
                 type="C-classification",
                 cost=myC, kernel="radial",
                 scale = FALSE)},
stop("Enter one of 'linear', 'poly.2', 'poly.3', 'radial'"))

x_valid <- valid[,1:2]
pred <- predict(model,x_valid)
t_true <- valid[,3]

# compute validation error for part 'i'
valid.error[i] <- sum(pred != t_true)/length(t_true)
}
# return average validation error
100*sum(valid.error)/length(valid.error)
}

```

Fit an SVM with linear kernel

In [28]: C <- 1000

```
(VA.error.linear <- train.svm.kCV ("linear", myC=C))
```

4

The procedure is to choose the model with the lowest CV error and then refit it with the whole learning data, then use it to predict the test set; we will do this at the end

Fit an SVM with quadratic kernel

In [29]: (VA.error.poly.2 <- train.svm.kCV ("poly.2", myC=C))

4

Fit an SVM with cubic kernel

In [30]: (VA.error.poly.3 <- train.svm.kCV ("poly.3", myC=C))

1

we get a series of decreasing CV errors ...
and finally an RBF Gaussian kernel

```
In [31]: (VA.error.RBF <- train.svm.kCV ("RBF", myC=C))
```

0.75

Now in a real scenario we should choose the model with the lowest CV error which in this case is the RBF (we get a very low CV error because this problem is easy for a SVM) so we choose RBF and C=1 and refit the model in the whole training set (no CV)

```
In [32]: model <- svm(dataset[,1:2],dataset[,3], type="C-classification", cost=C, kernel="radial"
```

and make it predict a test set:

let's generate the test data

```
In [33]: dataset.test <- make.sinusoidals (1000)
```

and have a look at it


```
In [34]: summary(dataset.test)
```

	x1	x2	target
Min.	:0.003142	Min. : -1.561915	-1:1000
1st Qu.	:1.573153	1st Qu.: -0.683925	1 :1000
Median	:2.357765	Median : -0.013291	
Mean	:2.357765	Mean : -0.008065	
3rd Qu.	:3.142378	3rd Qu.: 0.669875	
Max.	:4.712389	Max. : 1.661682	

```
In [35]: options(repr.plot.width=4, repr.plot.height=4)
```

```
par(mfrow=c(1,1))
```

```
plot(dataset.test$x1,dataset.test$x2,col=dataset.test$target)
```



```
In [36]: pred <- predict(model,dataset.test[,1:2])
         t_true <- dataset.test[,3]
```

```
         table(pred,t_true)
```

```
      t_true
pred  -1    1
  -1  996   0
   1    4 1000
```

compute testing error (in %)

```
In [37]: (sum(pred != t_true)/length(t_true))
```

```
0.002
```

In a real setting we should also optimize the value of C , again with CV; all this can be done very conveniently using `tune()` to do automatic grid-search (very much as we did in the last laboratory for `nnet()`)

other packages provide with heuristic methods to estimate the gamma in the RBF kernel (see below)

1.2 Playing with the SVM for regression and 1D data


Now we do regression; we have an extra parameter: the 'epsilon', which controls the width of the epsilon-insensitive tube (in feature space)

```
In [38]: A <- 20
```

a really nice-looking function


```
In [39]: x <- seq(-A,A,by=0.11)
         y <- sin(x)/x + rnorm(x,sd=0.03)

         plot(x,y,type="l")
```




With this choice of the 'epsilon', 'gamma' and C parameters, the SVM underfits the data (blue line)

```
In [40]: model1 <- svm (x,y,epsilon=0.01)
         plot(x,y,type="l")
         lines(x,predict(model1,x),col="blue")
```




With this choice of the 'epsilon', 'gamma' and C parameters, the SVM overfits the data (green line)

```
In [41]: model2 <- svm (x,y,epsilon=0.01,gamma=200, C=100)
          plot(x,y,type="l")
          lines(x,predict(model2,x),col="green")
```



With this choice of the 'epsilon', 'gamma' and C parameters, the SVM has a very decent fit (red line)

```
In [42]: model3 <- svm (x,y,epsilon=0.01,gamma=10)
         plot(x,y,type="l")
         lines(x,predict(model3,x),col="red")
```



the other nice package where the SVM is located is {kernlab}

```
In [43]: library(kernlab)
```

the `ksvm()` method has some nice features, as creation of user-defined kernels (not seen in this course) and automatic cross-validation (via the 'cross' parameter)